Scaling blockchain cryptocurrencies with leveled regions

Eduard Sanou
University of Southern California
Los Angeles, California
sanou@usc.edu

Borys Gurtovyi
University of Southern California
Los Angeles, California
gurtovyi@usc.edu

System	tx/sec	System	tx/sec
Bitcoin	3-4	Paypal	193
Ethereum	20	Visa	2,000
Bitcoin Cash	60		

Table 1: Transaction platforms and their throughput

ABSTRACT

Current blockchain based cryptocurrencies achieve very low transaction throughputs compared to widespread transaction systems like VISA or Mastercard. This makes current cryptocurrencies unfit to replace fiat currency. Based on the observation that most of the transactions[1] in a currency happen locally, and that users of a currency are used to expect lower confirmation delays on smaller distance transactions, we propose a way to scale a blockchain system. Our idea is based on a hierarchy of regions that allows transactions to happen locally first, and non-locally afterwards following periodic intervals that increase according to the distance.

1 INTRODUCTION

Current decentralized cryptocurrencies based on blockchains suffer from a scalability problem that makes them unfit to replace centralized payment systems such as VISA and PayPal: the transaction throughputs of cryptocurrencies is far from reaching that required for fiat currency transactions.

In this work, we propose an idea to achieve a higher scalability than current cryptocurrencies by creating a tree hierarchy of transaction regions at different levels, in which the transactions show different properties (in terms of cost and delay) depending on the lowest level ancestor in the tree of the sender and receiver.

We suggest using 4 levels in the hierarchy: *Global, Country, State, City*.

Our first proposal is for each lowest level to work on a fork of a global common blockchain, that would be merged with sibling regions (those under the same parent) periodically up in the hierarchy, with the periodicity being different at each level. The periodicity would be lower on higher level regions, and decrease for lower regions.

Transactions within the lowest level regions would be applied and verified quickly in a local blockchain whereas those traversing regions would only be confirmed by receivers once the forks of the corresponding regions are merged.

Our second proposal improves the efficiency of the first proposal by only requiring transactions that cross regions to be processed at higher levels in the hierarchy. Each region at each level would have a separate blockchain that would represent the underlying regions

---

[1]We wish to focus on transactions performed to exchange value rather than speculative transactions that happen in global currency trading

as a single *superwallet*, except for the local regions which would have a blockchain with regular wallets.

These designs would allow the transaction throughput within a single city to achieve a throughput close to what current blockchains achieve globally.

In other words, we are optimizing transaction delays by distance, which is a common expectation of fiat transactions, while allowing the global transaction throughput to increase significantly.

2 BACKGROUND

All existing blockchain protocols have a serious limitation: scalability. It is impossible to scale popular consensus systems such as Bitcoin because every single node on the network processes every transaction and maintains a copy of the entire state. The benefits of decentralization imply that every node is independent and has equal processing requirements.

With regular database system - the solution would be very easy. We could just add more servers - the number of servers is in direct ratio with the how system can scale. Adding more computing power may also help but in the decentralized world this means increasing the power of every node. We cannot force every participant to do so.

In fact, the blockchain actually gets weaker as more nodes are added to its network because of the inter-node latency that logarithmically increases with every additional node. Moreover, requirements for the node increase, and at some point risk of centralization appears (only some part of nodes are able to process the transaction). We end up with the trade off between throughput and decentralization, leading to low throughput in cryptocurrencies.

However this is not enough to meet the needs of fiat currency. Nowadays, Paypal processes 193 transactions per second and Visa 2000 (with a capability of 24,000). The majority of all transactions are local transaction (people do fiat currency transactions locally more often than to outside places). This fact pushed us to the idea of scaling blockchain by leveled regions.

3 SCALING WITH LEVELED REGIONS

The main idea behind our proposal is to scale a decentralized transaction system by removing the load nodes require to maintain the global data structure that is the blockchain. In order to remove this load, and thus allow an increase of transaction throughput,

we divide the work into a hierarchy, which for practical purposes is based on geographical regions. The following subsections will explain the different design choices of the system based on the Divide and Conquer approach used in algorithm design.

3.1 Hierarchical regions (*Divide*)

The decision to use geographical regions is arbitrary, and is motivated by the ease of understanding that the users will have upon the system, which can help with adoption. Other hierarchical divisions are possible with similar results; even dynamic regions that get constantly adjusted could provide efficiency benefits: for example, regions could be created such that all the regions in the same level have the same amount of users, or the same amount of computational power, or the same stake. Nevertheless, these more efficient divisions would make it harder for users to reason about which region they belong to: they are less intuitive.

For our design we will use political geographical regions divided into 4 levels: Global, Country, State, City.

The reasoning behind the numbers that identify the levels is the following: the global level is the root of the hierarchy and represents the top, so it will have the highest number; the rest of the levels are nested and have a decreasing decreasing number relative to their depth. We want our system to be flexible, and possible allow a finer level under a city (for example, big cities could be divided into neighborhoods). Following these constraints, we set the global level with the identifying number of 0, and decrement by one at each sub-level.

An important note on the divisions is that unlike other works [1] we don't require proof of location. Regions are made in order to divide the workload of transactions, with the idea of having local regions working on validating and accepting transactions happening on the local area; but this must be seen only as a convention to split the work (following the divide and conquer strategy). Should an attacker create identities that claim to be in a arbitrary geographical regions, the entire system security wouldn't be undermined. At most the efficiency of the entire system could be altered due to a possible unbalance of work between regions at the same level.

3.2 Local transactions (*Base case*)

Following the divide and conquer analogy, the lowest level regions (cities in our proposal) are the smaller partition of work. Transactions within a city (that is, transactions where the sender and receiver belong to the same city) are considered local, and are appended to a local blockchain, that for now, is independent of the other cities. This blockchain would work very similar to current Blockchain designs, with consensus achieved through Proof of Work, Proof of Stake, or any other proposed consensus mechanism. The key point here is that each city-level Blockchain will perform similarly to current Blockchain designs, and so, globally we are increasing the total throughput by several orders of magnitude. Moreover, since these Blockchains have a lower number of participants than global Blockchains, the block creation rate could be increased rendering a much lower transaction delay than current designs.
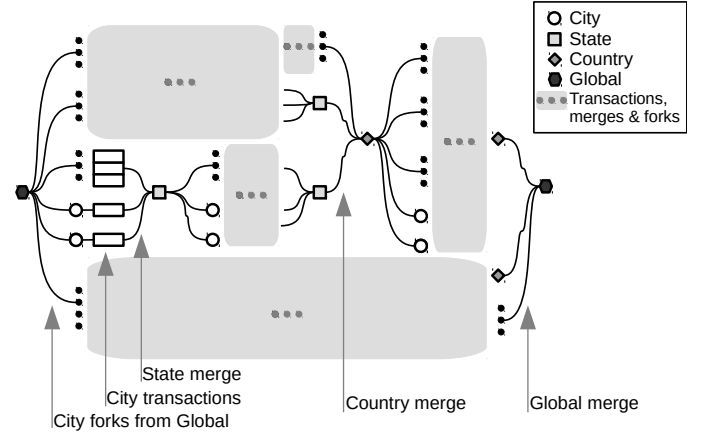

Figure 1: Work flow of forks and merges across the different levels of the hierarchy. Block transactions are only shown after the cities fork from the global chain for simplicity: block transactions are created at the local level after every merge-fork step, sown as gray dots in the figure.

Name	Level	periodicity
City	(-3)	10 seconds
State	(-2)	10 minutes
Country	(-1)	2 hours
Global	(0)	24 hours

Table 2: Geographical levels and their merging periodicity

Another possible design choice here is to use a permissioned blockchain for each city; considering that the nodes will be in the same community could help in the organization of such system.

3.3 Transacting across regions (*Conquer*)

If we stop here, what we would have is independent Blockchains for each city, each representing a different currency; but we are interested in allowing transactions across cities, across states and across countries.

3.3.1 Forks and merges. Our initial proposal is to start the system with a global Blockchain, common to all regions in all levels. This global Blockchain would be forked by the cities and periodically merged up in the hierarchy with the forks of the neighboring regions (at different intervals depending on the level) until reaching the global level, at which point the process would start again periodically. Figure 1 shows the work flow of this process.

Table 2 presents a proposal for interval times for merges at each level. In practice, these numbers would need to be tested with a simulation of transactions to guarantee that they are allow the merging process to happen without overlapping with the next iteration :

Transactions made by users would only be published on the local Blockchains, and the local community of nodes will be in
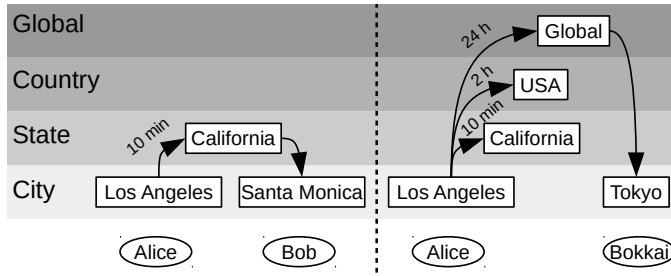
Figure 2: On the left, Alice (from *Los Angeles*) makes a transaction to Bob (in *Santa Monica*). Bob has to wait 10 minutes for the *Los Angeles* fork to be merged into the *California* merge. On the right, Alice (from *Los Angeles*) makes a transaction to Bokkai (in *Tokyo*). Bokkai has to wait 24 hours for the transaction to reach the *Global* merge.

charge of achieving consensus in order to guarantee correctness for local transactions but also to have an agreed validated list of transactions ready for the next merge with the rest of the cities in the same state. A rule that nodes would enforce is that they will only accept transactions (to be added to the city local blockchain) if the sender wallet belongs to the same city. Assuming that each city has reached consensus before the state merging point and that there are no *shadow*[2] cities, the state-level merge should proceed without conflicts.

Non-local transactions would be published on the local blockchain of the sender's wallet, but wouldn't be able to be confirmed by the receiver immediately. The receiver would have to wait until the sender's city fork gets merged with their own city fork, a delay that depends on the lowest common ancestor in the hierarchy. Figure 2 shows two examples of non-local transactions, and the required time for the receiver to be able to confirm them.

For our design, blockchain city-forks would be distinguished by a tag, so that all nodes claiming the same tag would work on the same blockchain. Having *shadow* cities can also be seen as not reaching consensus within a city, because it would imply disagreement under the same tag.

For wallets to belong to the same city, we introduce a special transaction that involves wallet creation: for a wallet to be used in regular transaction, a previous transaction which contains a city tag signed by the wallet private key must be published. Users traveling to a different city could create a new wallet belonging to the new city and transfer coins to it in order to perform local transactions in the new city; the traveler would just need to perform this operation before traveling to avoid any friction with the system. Like Bitcoin and other blockchain based cryptocurrencies, each user can have as many wallets as they desire; the only difference is that in our proposal, since a wallet can't be used without a transaction claiming a city where it belongs, there would probably be a small (transaction) cost associated to wallet ownership.

---

[2]We consider *shadow* cities the case where two set of nodes work on different blockchain forks but claim to be the same city.

3.3.2 Using different blockchains. The forking and merging strategy requires every node in the network to maintain a copy of the global blockchain, which contains a big number of transactions, thus increasing the storage requirements significantly. Moreover, the merging process requires a lot of bandwidth in order to synchronize the transactions between neighboring regions, with the worst case happening at the top level (where every country's fork merge into a global blockchain).

Considering that most of the transactions will happen locally, we could imagine a system in which nodes aren't required to store the transaction state of the entire system, as that information will be accessed less often. We could change the design so that local transactions stay in a local blockchain and are never merged up in the hierarchy: only transactions that need to cross regions move up and down the tree. This would reduce the bandwidth and storage requirements significantly. Furthermore, a mechanism could be engineered to charge higher transaction costs to further distance transactions, which would benefit users that do local transactions (and incentivize them to do so), and compensate nodes that work on non-local transactions.

So instead of working with a global blockchain that keeps forking and merging back following a tree; we could design the system to have a different blockchain for each region at each level, similar to what other works have proposed [1]. The question now is how transactions go up and down the hierarchy, and how nodes can verify the correctness of these vertical transactions.

Our proposal for this version of the system is to maintain the local blockchains the same way, and represent each region as a *superwallet* in the blockchain of the following upper level. For example, transactions between two cities would be seen as transactions between two *superwallets* in the corresponding state blockchain. We use the notion of *superwallets* to indicate the difference between regular local blockchain wallets: *superwallets* have the characteristic that they group many wallets underneath. Moreover, since transactions between *superwallets* represent transactions of lower level wallets, these transactions would need to carry extra information about the sender and receiver of each locally originated transaction. Figure 3 should help clarify this idea by sowing some examples of transactions in this model.

In this model, mining nodes could theoretically decide to work on the blockchains at different levels. Ideally, and to make things simpler, each node would work at each level (in the corresponding nested regions it belongs to). This would help linking a node to a local wallet in which the node would receive the mining rewards, which would need to be reflected in all levels to maintain correctness (that is, because the *superwallet* where the miner's wallet belongs would increase its coin value without receiving any transaction).

4 CHALLENGES/FUTURE WORK

* The merging process requires more computational work and bandwidth on higher levels. Not all nodes may be capable enough to verify merges at every level, Flexibility in the amount of contribution by node may be required.

* When a node performs a merge, it needs to verify the forks of the sibling regions for correctness. This could require a lot of computation resources. A possible optimization would be that each

State | Tx'(LA, SM, α) | α = Y + Z
California

City | Tx(A, B, X), Tx(C, SM:D, Y), Tx(E, SM:F, Z) | Tx(LA:C, D, Y), Tx(LA:E, F, Z)
Los Angeles | Santa Monica

Figure 3: The *Los Angeles* block contains three transactions, out of which one is local and two have recipients in *Santa Monica*. Non-local transactions go up into the hierarchy to the State level. In the California block only one transaction is shown, which corresponds to the non-local transactions originated at *Los Angeles,* grouped under a single superwallet. The *Santa Monica* block takes the transaction received by the *Santa Monica* supwerallet and decomposes them into individual transactions.

merge outputs a compressed form of the result that makes verification at the next level faster.

* At the lowest level, each region would have a small number of nodes. This opens the possibility of an attacker disrupting the region by registering nodes in it with high validation weight (either computational power or stake). Such disruption could be a denial of service or a denial of selected transactions.

* A way to make make sure that all regions have the same computing power could be useful to solve this.

* A different approach would consist on restricting how easily new nodes with high computational power can join or switch regions.

* At each level, the merging time between forks should be faster than the periodicity of merges.

5 RELATED WORK

5.1 Unchain Your Blockchain

This is a proposal [1] that tries to solve the current scalability issues of public blockchain based technologies. Very similar to our approach, the authors design a hierarchical division based on geographical regions. Each region at each level would have a different and independent blockchain such that local transactions happen in the local blockchain associated with the geographical area. Farther transactions require multiple transactions to move up and down in the hierarchy to pass through the lowest common region; increasing the cost and delay compared to local transactions.

The authors complement this design with the idea of Proof of Location, which is used by miners to maintain consensus of the blockchains in a fair manner and without the power consumption associated with Proof of Work. This design requires an infrastructure of trusted corroborators uniformly spread over the geographical area where the cryptocurrency is deployed. This corroborators consist of antennas that have the capability of issuing certificates to parties after they have verified their proximity (by analyzing the communication latency). The algorithm to select which miner creates the next block in the blockchain consist of the publication of a random corroborator antenna, so that the miner who can prove

to be the closest (in distance, or latency) to the antenna wins (and gets to create the block). Location certificates can also improve the security from the regular users point of view by means of geofencing their wallet, such that the coins in it can only be spent from a particular geographical area. The authors devise a business model behind the corroborators by charging a fee per issued certificate, which would increase with the distance to deter miners from getting too many of them.

The description of the system doesn't specify how correctness is verified when crossing levels. In single blockchain systems, each transaction is verified by asserting that the sender has enough coins through inspection of the entire transaction history. It's not clear how a blockchain could verify that the sender in a transaction from another level has enough coins, without having access to different level blockchain.

In contrast to our proposal, the above system requires the deployment of physical infrastructure from the corroborators, as well as trusting them in order to function with security guarantees. Our idea doesn't require any proof of location, and doesn't offer any advantage to an attacker that claims to be in a different location than their geophysical location.

5.2 Payment Channels

Payment channels are a proposal to improve the scalability of transactions by allowing payments to happen out of the blockchain, thus increasing the total transaction throughput and not requiring miners to process every transaction. Rather than creating a hierarchy of transactions, this solution creates a single off-chain protocol in which a group of parties can do an unlimited number of transactions, which is connected to the blockchain through only two transactions. The first blockchain transaction is a payment that the parties make to a channel wallet and the second transaction is a payment from such channel wallet back to the participants according to the final balance of each party in the off-chain transaction state.

Whereas initial proposals like the Lightning Network [6] or the work of Christian Decker et. al. [2] offer solutions for just scaling transaction throughput by allowing off-chain transactions, other works like TumbleBit [3] and the work of Giulio Malavolta et. al. [4] expand the idea by supporting anonymous payments thus increasing the privacy of the systems.

Our proposal could be seen as an extension of payment channels, where the channels are blockchains themselves and exist in a hierarchical tree based on geographic regions.

5.3 Plasma Project

The authors of Ethereum have proposed the Plasma Project [5] which extends the idea of off-chain transactions with smart contracts. The proposal involves a tree-like structure where a root blockchain has several child blockchains that handle most of the computation and transaction load, only using the root blockchain to store small commitments to guarantee the security of the nested transactions.

Their system still requires nodes to verify transactions of children blockchains, with the caveat that nodes can choose to only

work on children blockchains of interest (those that handle transactions or smart contracts that the node cares about). The data structure used in the root chain to secure transactions in children blockchains is the Merkle tree.

Our idea shares some concepts with the Plasma Project, although we focus on monetary transactions only. In contrast to Plasma, our hierarchical tree would be based on geographical regions.

5.4 Sharding

Ethereum presented idea of sharing in their own whitepaper. The point is the blockchain can have at most 2 to of these 3 properties: Decentralization, Scalability, Security. Currently, transactions on the EVM are not parallelizable, and every transaction is executed in sequence globally.

Sharing is splitting all nodes to a bunch of portions called shards. Every shard contains its own independent state and transactions history. Every shard contains special type of nodes called collators for storing information about particular shard. This information has what shard is applies to, state of the shard before/after all transaction are applied, digital signature of > 2/3 all collators on shard which make this collation legit.

Second important type of node is supernode which is in this case responsible for taking all collations across the shard and put them into single block. This block will be added to blockchain.

The new block is valid if and only if all transactions in all collations are valid, the state of the collations is same as the current state of the collations before the transactions, the state of collations after transactions is same as specified and 2/3 collators signed collations. The main question now is how to send transactions across the shards. For example A (shard 1) sends a 1 ETH to B (shard 2).

A transaction is sent to Shard 1 $\rightarrow$ A's balance is reduced by 1 ETH.

A receipt is created. It is not being stored in a state yet, but in the Merkle root.

Merkle receipt with the transaction is sent to Shard 2. Balance of B is increased by 1 ETH.

The biggest problem is a Single-Shard takeover attack. It is enough now for attacker to take over 1 shard to start submitting invalid collations. The Ethereum Sharding FAQ suggests random sampling of collators on each shard.

6 CONCLUSIONS

REFERENCES

[1] Tamraparni Dasu, Yaron Kanza, and Divesh Srivastava. 2018. Unchain your blockchain. In *Proc. Symposium on Foundations and Applications of Blockchain*, Vol. 1. 16–23.
[2] Christian Decker and Roger Wattenhofer. 2015. A fast and scalable payment network with bitcoin duplex micropayment channels. In *Symposium on Self-Stabilizing Systems*. Springer, 3–18.
[3] Ethan Heilman, Leen Alshenibr, Foteini Baldimtsi, Alessandra Scafuro, and Sharon Goldberg. 2017. Tumblebit: An untrusted bitcoin-compatible anonymous payment hub. *Proceedings of NDSS 2017* (2017).
[4] Giulio Malavolta, Pedro Moreno-Sanchez, Aniket Kate, Matteo Maffei, and Srivatsan Ravi. 2017. Concurrency and privacy with payment-channel networks. In *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security*. ACM, 455–471.
[5] Joseph Poon and Vitalik Buterin. 2017. Plasma: Scalable Autonomous Smart Contracts. *White paper* (2017).
[6] Joseph Poon and Thaddeus Dryja. 2016. The bitcoin lightning network: Scalable off-chain instant payments. *draft version 0.5* 9 (2016), 14.